

---

# Student Dropout Prediction: A Machine Learning Analysis

DSAA2011 (L01) Course Project — HKUST(GZ) 2026 Spring

---

Xingjian LI   Congren HANG   Hao LUO   Ruoxi WANG

## Abstract

We analyze the UCI Student Dropout and Academic Success dataset through preprocessing, t-SNE visualization, clustering, supervised classification, model evaluation, and open-ended exploration. The engineered top-40 feature space is dominated by academic performance signals such as semester approval rates and grades. Unsupervised analysis shows that Enrolled students lie between Dropout and Graduate rather than forming a clean natural cluster. Logistic Regression is the strongest mandatory baseline. In open-ended exploration, validation-tuned class-bias adjustment for LR achieves the best held-out test macro F1 of 0.7228, improving the Enrolled decision boundary while remaining only borderline significant over plain LR.

## 1 Introduction

The Student Dropout and Academic Success dataset from the UCI Machine Learning Repository [1] contains 4,424 instances of students enrolled in various undergraduate programs at a Portuguese higher education institution. Each instance is described by 36 features covering demographics, socioeconomic background, and academic performance across the first and second semesters.

The report follows an evidence-first structure: construct an engineered top-40 feature space, use t-SNE and clustering to test whether outcomes form natural groups, then evaluate supervised decision boundaries. Because Enrolled students remain structurally intermediate between Dropout and Graduate, the open-ended section focuses on class-boundary refinement rather than only adding model complexity.

## 2 Mandatory Tasks

### 2.1 Data Preprocessing

The dataset contains 36 raw features that are categorized into four types: *nominal* (8 features), *binary* (8 features), *ordinal* (13 features), and *continuous* (7 features). It also includes a three-class target variable: Dropout (32.1%), Enrolled (18.0%), and Graduate (49.9%). We performed a comprehensive data preprocessing pipeline as outlined below.

**Missing values and outliers.** The dataset is clean with no missing values, as confirmed by a complete null-value check.

**EDA for nominal features.** We inspected the distribution and top categories of nominal features. Except for feature *Course*, most nominal features are heavily skewed and the top four categories can cover over 80% of the data. This informed our decision to apply out-of-fold (OOF) target encoding for these features, as one-hot encoding would lead to high dimensionality and sparsity.

**EDA for numerical correlations.** We sorted numerical features by absolute Pearson correlation with the target and visualized the top 15 with  $|r| \geq 0.09$  (Figure 1). Curricular-unit approval and grade features show the strongest signal, while economic and demographic variables provide weaker but still useful complementary information.




Figure 1: Top-15 feature correlations with the target variable.

**EDA for important features.** We visualized four key curricular-unit features as histograms with trendlines (Figure 2). Higher values are associated with Graduate, lower values with Dropout, and Enrolled students remain intermediate.




Figure 2: Distributions of top curricular unit features across target classes.

**Feature engineering and selection.** We created 25 domain-specific features from the 36 raw features, yielding 61 candidates.

**Data splitting and target encoding.** The dataset was split into stratified train/validation/test sets (70/15/15). OOF target encoding was fit on training folds for nominal features and then applied to validation and test sets, preventing leakage while capturing categorical signal.

**Feature standardization, selection, and output.** A Random Forest selected the top 40 target-encoded features by importance. Non-binary selected features were standardized using training-set statistics only; binary indicators remained 0/1, and the fitted scaler was reused for validation and test sets.

## 2.2 Data Visualization (t-SNE)

We applied t-SNE to the engineered feature dataset to project the high-dimensional feature space into lower dimensions for visualization.

A systematic pattern analysis across four perplexity values 10, 20, 30, and 50 highlights the algorithm’s sensitivity to local versus global structure, as shown in Figure 3.

At a low perplexity ( $p = 10$ ), the projection is severely fragmented. Higher perplexity values ( $p = 20, 30, 50$ ) produce similar stable global structures and the relative positions of the three classes remain consistent. However, high perplexity tends to over-smooth the local structure and concentrate points into denser clusters. Generally, perplexity 30 achieves the optimal balance.

The 2D projection reveals that Graduate and Dropout students form relatively compact clusters, while Enrolled students scattered between both groups and mostly overlap with the Graduate cluster, indicating that Enrolled students do not form a distinct cluster but rather occupy an intermediate region in the feature space. This pattern is consistent across all perplexity values, confirming that Enrolled students have transitional academic status. The overlapping nature of the Enrolled class suggests that it may be more challenging to classify them accurately based on the engineered features or clustering alone.

Furthermore, extending the projection from 2D to 3D (Figure 4) provides additional spatial volume, further reducing the overlap between the distinct Graduate and Dropout clusters.




Figure 3: t-SNE 2D projections at perplexity 10, 20, 30, 50.

Figure 4: t-SNE 3D projection at perplexity 30.

## 2.3 Clustering Analysis

We applied three clustering algorithms to the engineered top-40 feature space: K-Means, Agglomerative Hierarchical Clustering, and Gaussian Mixture Models (GMM). DBSCAN and OPTICS were explored but failed: DBSCAN classified 100% of points as noise across all  $\epsilon$  values (0.3–1.5), and OPTICS with  $\xi$ -based extraction consistently returned only one cluster. Both failures stem from the curse of dimensionality — in 40 dimensions all pairwise distances become similar (coefficient

of variation = 0.32), making density-based methods unable to identify dense regions. PCA was used for visualization as it is stable and distance-consistent.

**K-Means K-Value Exploration:** We tested K=3,4,5,6. Although K=3 matches the number of target classes, K=4 nearly doubled the ARI from 0.171 to 0.313 (Table 1), suggesting that target-encoded features benefit from finer cluster granularity. K=4 was selected for detailed analysis.

Table 1: K-Means K-value exploration on engineered features.

| K | Silhouette | CH Index | DB Index | ARI           | NMI           |
|---|------------|----------|----------|---------------|---------------|
| 3 | 0.2741     | 839.73   | 1.5640   | 0.1708        | 0.1757        |
| 4 | 0.1704     | 719.87   | 2.0741   | <b>0.3129</b> | <b>0.2299</b> |
| 5 | 0.1690     | 625.90   | 1.9093   | 0.3052        | 0.2213        |
| 6 | 0.1243     | 566.57   | 1.9614   | 0.2061        | 0.1861        |

**Hierarchical Clustering:** We first compared four linkage methods with K=3 (Table 2). Ward linkage achieved the highest ARI (0.175) by far, while single and average linkages produced degenerate partitions (ARI  $\approx$  0). K-value exploration with Ward linkage confirmed K=3 as optimal; larger K values did not improve ARI. The dendrogram (Figure 5) shows a relatively flat merge structure.




Figure 5: Hierarchical clustering dendrogram (Ward linkage).

Table 2: Linkage methods comparison (K=3).

| Method   | Silhouette | CH Index | DB Index | ARI           |
|----------|------------|----------|----------|---------------|
| Ward     | 0.2436     | 725.18   | 1.7059   | <b>0.1750</b> |
| Complete | 0.2256     | 214.30   | 1.6389   | -0.0095       |
| Average  | 0.5374     | 22.10    | 0.7102   | 0.0009        |
| Single   | 0.5290     | 7.49     | 0.3460   | 0.0002        |

**GMM K-Value and Covariance Exploration:** We tested K=3,4,5,6 combined with four covariance types (full, tied, diag, spherical). The tied covariance structure performed best, and K=4 achieved the overall best ARI of 0.334 (Table 3). This confirms that allowing all clusters to share the same covariance shape reduces overfitting in the 40-dimensional space while still capturing cluster differences through means.

Table 3: GMM covariance and K exploration (best configurations).

| K | Covariance | Silhouette | CH Index | DB Index | ARI           | NMI           |
|---|------------|------------|----------|----------|---------------|---------------|
| 3 | tied       | 0.1510     | 741.35   | 2.4022   | 0.3324        | 0.2578        |
| 3 | diag       | 0.1670     | 712.43   | 2.7407   | 0.3334        | 0.2505        |
| 4 | tied       | 0.1423     | 559.91   | 2.2580   | <b>0.3336</b> | <b>0.2734</b> |
| 4 | diag       | 0.1627     | 664.36   | 2.3042   | 0.3213        | 0.2348        |

**Overall Algorithm Comparison:** Table 4 summarizes all three algorithms with their optimal configurations. GMM (K=4, tied) leads in both ARI (0.334) and NMI (0.273), substantially outperforming K-Means and Hierarchical clustering. Only GMM produces dropout-dominant clusters (Table 5); K-Means K=4 and Hierarchical produce graduate-dominant clusters, indicating that soft probabilistic boundaries better capture the structural overlap between Dropout and Graduate. Figure 9 visualizes the three partitions on a 2D PCA projection.

Table 4: Clustering algorithm performance comparison.

| Algorithm           | K | Silhouette | CH Index | DB Index | ARI           | NMI           |
|---------------------|---|------------|----------|----------|---------------|---------------|
| K-Means             | 4 | 0.1704     | 719.87   | 2.0741   | 0.3129        | 0.2299        |
| Hierarchical (Ward) | 3 | 0.2436     | 725.18   | 1.7059   | 0.1750        | 0.1609        |
| GMM                 | 4 | 0.1423     | 559.91   | 2.2580   | <b>0.3336</b> | <b>0.2734</b> |

Table 5: Dominant true class per cluster for each algorithm.

| Algorithm          | Cluster 0        | Cluster 1       | Cluster 2        | Cluster 3       |
|--------------------|------------------|-----------------|------------------|-----------------|
| K-Means (K=4)      | Graduate (76.2%) | Dropout (52.6%) | Graduate (66.7%) | Dropout (80.3%) |
| Hierarchical (K=3) | Graduate (57.3%) | Dropout (54.5%) | Graduate (54.6%) | —               |
| GMM (K=4, tied)    | Graduate (76.2%) | Dropout (59.6%) | Graduate (66.7%) | Dropout (80.3%) |




Figure 6: K-Means (K=4)




Figure 7: Hierarchical (Ward, K=3)



Figure 8: GMM (K=4, tied)

Figure 9: Cluster visualization on 2D PCA projections. red - Dropout, orange - Enrolled, green - Graduate.

**Best clustering: GMM (K=4, tied covariance)** achieves the best ARI of 0.334 and NMI of 0.273, confirming that the shared covariance structure reduces overfitting in the high-dimensional target-encoded space. The modest ARI still reflects structural class overlap (Enrolled students lie between Dropout and Graduate), but GMM’s soft assignments capture more label-aligned structure than centroid-based methods.

## 2.4 Prediction: Training and Testing

The clustering results suggest that the three student outcomes are not naturally separated in the engineered feature space. We therefore train supervised classifiers on the top-40 engineered features to test whether label-informed decision boundaries can better distinguish Dropout, Enrolled, and Graduate students.

**Data preparation.** We used the engineered feature dataset (40 features after OOF target encoding and top-40 selection), split into train (3,096), validation (664), and test (664) sets. The class distribution in the training set is: Dropout 32.1%, Enrolled 17.9%, Graduate 49.9%.

Decision Tree (DT), Logistic Regression (LR), and RBF-SVM were selected because they represent complementary assumptions: non-linear splits, interpretable linear boundaries, and kernel-based nonlinear boundaries.

We trained each classifier using default configurations from `prediction.py`: Decision Tree (`max_depth=10`, `class_weight='balanced'`), Logistic Regression (`solver='saga'`, `class_weight='balanced'`), and SVM with RBF kernel (`class_weight='balanced'`). All models were evaluated on training, validation, test, and entire sets.

**Multi-set evaluation.** Table 6 reports accuracy and macro F1 across all evaluation sets. DT has the largest train-test gap, indicating overfitting. LR has the strongest held-out test macro F1 among the mandatory baselines and remains stable across train, validation, and test splits. Default RBF-SVM improves substantially after the corrected scaling step, but LR remains the strongest mandatory baseline by held-out macro F1.

Table 6: Baseline model performance (macro F1 across splits and test accuracy).

| Model         | Train F1 | Val F1 | Test F1 | Test Acc. |
|---------------|----------|--------|---------|-----------|
| Decision Tree | 0.8876   | 0.6709 | 0.6752  | 0.7123    |
| Logistic Reg. | 0.7163   | 0.7076 | 0.7076  | 0.7410    |
| SVM (RBF)     | 0.7752   | 0.6937 | 0.6955  | 0.7319    |

Logistic Regression — Test Set Confusion Matrix

|            |                 |               |                |
|------------|-----------------|---------------|----------------|
| True Label | Dropout         | Enrolled      | Graduate       |
|            | 150<br>(70.4%)  | 50<br>(23.5%) | 13<br>(6.1%)   |
|            | 15<br>(12.6%)   | 81<br>(68.1%) | 23<br>(19.3%)  |
| Graduate   | 14<br>(4.2%)    | 57<br>(17.2%) | 261<br>(78.6%) |
|            | Dropout         | Enrolled      | Graduate       |
|            | Predicted Label |               |                |

Figure 10: Logistic Regression confusion matrix on the test set.

The per-class analysis shows that all models achieve the highest recall for Graduate and the lowest for Enrolled. This is consistent with the class imbalance and the transitional nature of enrolled students, who share feature characteristics with both Dropout and Graduate.

**Error analysis.** The confusion matrix (Figure 10) reveals the dominant error pattern: Graduate students are correctly identified with the highest recall (78.6%), while Enrolled students show the lowest (68.1%), with 12.6% misclassified as Dropout and 19.3% as Graduate. This confirms that Enrolled does not occupy a distinct decision region. The pattern aligns with the t-SNE visualization (Section 2.2) and clustering results (Section 2.3), where Enrolled points lie in the overlapping zone between the two more separable outcomes. The structural overlap between classes, beyond class imbalance alone, is a major source of classification difficulty.

## 2.5 Evaluation and Model Choice

We evaluate the three mandatory baselines using held-out precision/recall/F1, One-vs-Rest ROC-AUC, cross-validation, validation curves, learning curves, and GridSearchCV tuning. The overall final model choice is deferred to the open-ended section.

**Detailed metrics.** Table 7 reports the full set of evaluation metrics on the test set. LR leads among the mandatory baselines by macro F1, while SVM becomes competitive after corrected scaling but still trails LR.

Table 7: Detailed per-model metrics on the test set.

| Model         | Acc.   | Prec. (macro) | Recall (macro) | F1 (macro) |
|---------------|--------|---------------|----------------|------------|
| Decision Tree | 0.7123 | 0.6859        | 0.6908         | 0.6752     |
| Logistic Reg. | 0.7410 | 0.7159        | 0.7237         | 0.7076     |
| SVM (RBF)     | 0.7319 | 0.7054        | 0.7096         | 0.6955     |

Table 8: Per-class and Macro AUC.

| Model | Dropout      | Enrolled     | Graduate     | Macro        |
|-------|--------------|--------------|--------------|--------------|
| DT    | 0.797        | 0.702        | 0.849        | 0.783        |
| LR    | <b>0.910</b> | <b>0.822</b> | 0.932        | <b>0.888</b> |
| SVM   | 0.909        | 0.821        | <b>0.932</b> | 0.887        |

**ROC curves and AUC.** Figure 11 compares macro-average ROC curves. LR achieves the highest macro AUC (0.888), closely followed by SVM (0.887) and DT (0.783).

**Cross-validation.** 5-fold CV on the combined train+val set (3,760 samples) confirms LR as the most stable mandatory baseline: LR CV F1 =  $0.7029 \pm 0.0064$ , DT =  $0.6511 \pm 0.0144$ , SVM =  $0.6955 \pm 0.0063$ .

**Validation curve analysis** reveals model-specific sensitivity to hyperparameters. Figure 12 shows the two most informative curves: DT max\_depth exhibits classic overfitting (training score rises while validation drops), while SVM C shows that increasing regularization flexibility steadily improves both training and validation performance.

LR is relatively insensitive to its regularization parameter C, suggesting the linear assumption is the primary bottleneck. SVM gamma (not shown) causes severe overfitting at high values; the optimal C and gamma must be tuned jointly.




Figure 11: Macro-average ROC curves for all baseline models.




Figure 12: Validation curves: DT max\_depth (left) shows overfitting; SVM C (right) shows improvement with larger C.

**Learning curves** confirm that DT persistently overfits, LR has nearly converged, and SVM could benefit from more data.

**Hyperparameter tuning via GridSearchCV** produced the best configurations:

**DT:** max\_depth=5, min\_samples\_split=2, min\_samples\_leaf=1 (CV F1=0.6839)  
**LR:** C=10, penalty='l2', solver='saga' (CV F1=0.7051)  
**SVM:** C=10, gamma=0.01 (CV F1=0.7039)

Table 9: Baseline vs. tuned (F1 macro).

| Model | Baseline | Tuned  | $\Delta$ |
|-------|----------|--------|----------|
| DT    | 0.6752   | 0.6731 | -0.0021  |
| LR    | 0.7076   | 0.7074 | -0.0002  |
| SVM   | 0.6955   | 0.7033 | +0.0078  |

SVM benefits most from tuning, improving held-out macro F1 from 0.6955 to 0.7033. DT and LR do not improve on the test set, suggesting that LR's limitation is its linear boundary rather than regularization strength. Enrolled remains the hardest class across all configurations.

**Baseline summary.** Within the mandatory baselines, LR is strongest and most stable. DT overfits, SVM benefits from tuning but remains below LR, and Enrolled remains the hardest class because it is both underrepresented and structurally intermediate. These findings motivate open-ended exploration beyond baseline deployment.

### 3 Open-ended Exploration

#### 3.1 Ensemble and Decision-boundary Refinement

The corrected top-40 feature space changes the open-ended question: the baseline bottleneck is no longer simply weak model capacity, but the boundary around the Enrolled class. We therefore evaluated tree ensembles (RF, GB, ExtraTrees) and a validation-tuned LR decision rule. The latter keeps the LR probability model fixed but tunes additive class biases on the validation set to maximize macro F1, then applies the selected biases to the held-out test set.

Table 10: Open-ended model and decision-rule performance on the test set.

| Model                      | Acc.          | Prec. (macro) | Recall (macro) | F1 (macro)    | F1 (wtd)      |
|----------------------------|---------------|---------------|----------------|---------------|---------------|
| RF (default)               | 0.7726        | 0.7163        | 0.6867         | 0.6942        | 0.7600        |
| GradientBoosting (default) | 0.7545        | 0.6883        | 0.6821         | 0.6848        | 0.7516        |
| RF (tuned)                 | 0.7470        | 0.7120        | 0.7133         | 0.7064        | 0.7554        |
| GradientBoosting (tuned)   | <b>0.7726</b> | 0.7131        | 0.6965         | 0.7027        | 0.7657        |
| ExtraTrees                 | 0.7666        | 0.7174        | 0.7141         | 0.7143        | 0.7689        |
| LR + class-bias tuning     | 0.7651        | <b>0.7303</b> | <b>0.7290</b>  | <b>0.7228</b> | <b>0.7725</b> |




Figure 13: ROC curves for tuned RF, tuned GB, and ExtraTrees.

Ensemble gains are useful but modest: ExtraTrees reaches macro F1 = 0.7143, above RF and GB, while the validation-tuned LR rule gives the largest held-out improvement (0.7228). This supports the interpretation that the remaining error is mainly a class-boundary problem around Enrolled. The selected validation biases [-0.8, -0.8, -0.4] adjust LR's relative class thresholds and improve macro balance.



### 3.2 Class Imbalance Handling

Table 11 compares per-class Enrolled F1 with and without balanced class weighting across the three mandatory baseline models after the corrected standardization step.

Table 11: Effect of class weighting on Enrolled F1 score.

| Model         | Enrolled F1 (default) | Enrolled F1 (balanced) | $\Delta$ |
|---------------|-----------------------|------------------------|----------|
| Decision Tree | 0.4034                | 0.5016                 | +0.0982  |
| Logistic Reg. | 0.4476                | 0.5277                 | +0.0801  |
| SVM (RBF)     | 0.4887                | 0.5180                 | +0.0293  |

Balanced weighting still improves Enrolled recognition, especially for DT and LR. SVM no longer collapses after scaling correction, so the previous SVM failure was partly preprocessing-related. Remaining Enrolled errors reflect both minority-class pressure and structural overlap.

### 3.3 Cross-Validation Comparison

A unified 5-fold cross-validation (Figure 14) compared six model classes on the training set:

Table 12: 5-fold CV F1 (macro) results.

| Model      | CV F1 (Mean $\pm$ Std) | Note                          |
|------------|------------------------|-------------------------------|
| ExtraTrees | 0.7170 $\pm$ 0.0171    | Highest mean                  |
| RF         | 0.7101 $\pm$ 0.0208    | Strong ensemble               |
| GB         | 0.7051 $\pm$ 0.0093    | Stable ensemble               |
| LR         | 0.7021 $\pm$ 0.0079    | Strong interpretable baseline |
| SVM        | 0.6938 $\pm$ 0.0153    | Improved after scaling        |
| DT         | 0.6507 $\pm$ 0.0118    | Overfitting risk              |




Figure 14: 5-fold CV comparison across model classes.

ExtraTrees has the highest mean CV macro F1, while LR remains close and more interpretable. CV shows that model-family gains are modest; separately, the validation-tuned LR decision rule gives the largest held-out test gain.

### 3.4 Model Interpretation and Feature Ablation

Feature importance scores from the tuned Random Forest (Figure 15) confirm that academic performance metrics dominate prediction. The top-10 features are led by `overall_approval_rate` and `approval_rate_2nd`, followed by curricular unit metrics and semester grades.




Figure 15: Top-10 feature importances from the tuned Random Forest.

**Feature group ablation.** Academic features alone account for most predictive signal (F1 = 0.669), while socioeconomic features alone are weaker (F1 = 0.561). Combining both groups yields the best and most stable performance (F1 = 0.714), confirming complementary information beyond academic records.

Table 13: Feature group ablation (5-fold CV, tuned RF).

| Feature Group      | $n$ | CV F1 (Mean $\pm$ Std) |
|--------------------|-----|------------------------|
| Academic only      | 20  | 0.6692 $\pm$ 0.0116    |
| Socioeconomic only | 20  | 0.5607 $\pm$ 0.0245    |
| All top-40         | 40  | 0.7142 $\pm$ 0.0073    |



**Feature count and interaction experiments.** Feature count experiments (10/20/30/40 features, 5-fold CV) show that top-40 is best ( $F1 = 0.700$ ), while top-10 drops to 0.655. Polynomial interactions expanded 40 to 85 dimensions but did not improve CV  $F1$ , so we retain the original top-40 space.

### 3.5 Comprehensive Model Comparison

Table 14 presents all test-set model configurations. Using held-out macro  $F1$  as the criterion, LR with validation-tuned class biases is the final selected decision rule (0.7228), because it directly targets the observed Enrolled boundary problem.

**Per-class analysis.** Table 15 compares LR+bias with plain LR, both trained on train+validation data. The tuned rule improves Enrolled  $F1$  from 0.528 to 0.538 and Graduate  $F1$  from 0.829 to 0.854 while preserving Dropout performance.

Table 14: Comprehensive comparison of all models on the test set.

| Model                      | Accuracy      | $F1$ (macro)  | $F1$ (weighted) |
|----------------------------|---------------|---------------|-----------------|
| Decision Tree              | 0.7123        | 0.6752        | 0.7263          |
| Logistic Regression        | 0.7410        | 0.7076        | 0.7550          |
| SVM (RBF)                  | 0.7319        | 0.6955        | 0.7448          |
| RF (default)               | 0.7726        | 0.6942        | 0.7600          |
| GradientBoosting (default) | 0.7545        | 0.6848        | 0.7516          |
| RF (tuned)                 | 0.7470        | 0.7064        | 0.7554          |
| GradientBoosting (tuned)   | <b>0.7726</b> | 0.7027        | 0.7657          |
| ExtraTrees                 | 0.7666        | 0.7143        | 0.7689          |
| LR + class-bias tuning     | 0.7651        | <b>0.7228</b> | <b>0.7725</b>   |

Table 15: Per-class  $F1$  comparison: LR+bias vs. plain LR.

| Class    | LR+bias $F1$ | LR $F1$ | $\Delta$ |
|----------|--------------|---------|----------|
| Dropout  | 0.777        | 0.774   | +0.004   |
| Enrolled | 0.538        | 0.528   | +0.010   |
| Graduate | 0.854        | 0.829   | +0.025   |

**Statistical significance.** Bootstrap analysis (10,000 test-set resamples) gives a 95% CI for LR+bias – LR of  $[-0.0026, +0.0278]$  with  $p = 0.051$ . The interval narrowly crosses zero, so the gain is borderline rather than definitively significant. Still, the direction and error pattern support targeted decision-boundary refinement.

## 4 Conclusion

**Key findings:** The engineered top-40 feature space contains strong predictive signal, with academic metrics providing the primary contribution (academic-only  $F1 = 0.669$ ; all top-40  $F1 = 0.714$ ). t-SNE and clustering show that the outcomes are not naturally separated because Enrolled students lie between Dropout and Graduate. Among mandatory baselines, LR is strongest and most stable; DT overfits and SVM improves once the final top-40 space is correctly standardized. In open-ended exploration, LR with validation-tuned class biases achieves the best held-out test macro  $F1$  of 0.7228 versus 0.7102 for plain LR, though the bootstrap result is borderline ( $p = 0.051$ ). Enrolled remains the primary challenge because it is both a minority class and a structurally transitional outcome.

**Limitations:** The borderline gap between LR+bias and plain LR means the final result should not be treated as a definitive deployment advantage without external validation. The t-SNE decision-boundary plots are 2D approximations, and clustering struggled because discriminative features do not correspond to natural density-based clusters.

**Future work:** Future work should focus on better modeling of Enrolled students, external validation, and temporal academic trajectories for early-warning deployment.

## References

- [1] Realinho, V., et al. (2022). Student Dropout and Academic Success. UCI Machine Learning Repository. <https://doi.org/10.24432/C5MC89>
- [2] Pedregosa, F., et al. (2011). Scikit-learn: Machine Learning in Python. JMLR, 12, pp. 2825–2830.
- [3] Van der Maaten, L. and Hinton, G. (2008). Visualizing Data using t-SNE. JMLR, 9, pp. 2579–2605.
- [4] Claude by Anthropic (Claude Opus 4) was used for code debugging, visualization code generation, and report language polishing, accounting for less than 30% of total project effort. All core ML implementations are original student work.

## Credit

| Member       | Contribution                                               |
|--------------|------------------------------------------------------------|
| Hao LUO      | [Data preprocessing, feature engineering, report writing]  |
| Congren HANG | [Clustering analysis, t-SNE visualization, report writing] |
| Xingjian LI  | [Model training, evaluation, report writing]               |
| Ruoxi WANG   | [Open-ended exploration, presentation]                     |